

RepositoryCodes-Final1

October 7, 2019

1 Introduction

Body fat percentage is total mass of fat divided by total body mass, it is a important factor to show a person's health level, too high or too low will caused diseases. There are some body fat calculators online, however, there are too many variables. In this project, we try to find a simple, accurate and robust model to predict boday fat Percentage which obeys the 'Rule of Thumb'.

```
[1]: library(car)
      library(leaps)
      library(faraway)
      library(glmnet)
```

```
Loading required package: carData
Registered S3 methods overwritten by 'lme4':
  method                      from
  cooks.distance.influence.merMod car
  influence.merMod             car
  dfbeta.influence.merMod      car
  dfbetas.influence.merMod     car
```

```
Attaching package: 'faraway'
```

```
The following objects are masked from 'package:car':
```

```
  logit, vif
```

```
Loading required package: Matrix
Loading required package: foreach
Loaded glmnet 2.0-18
```

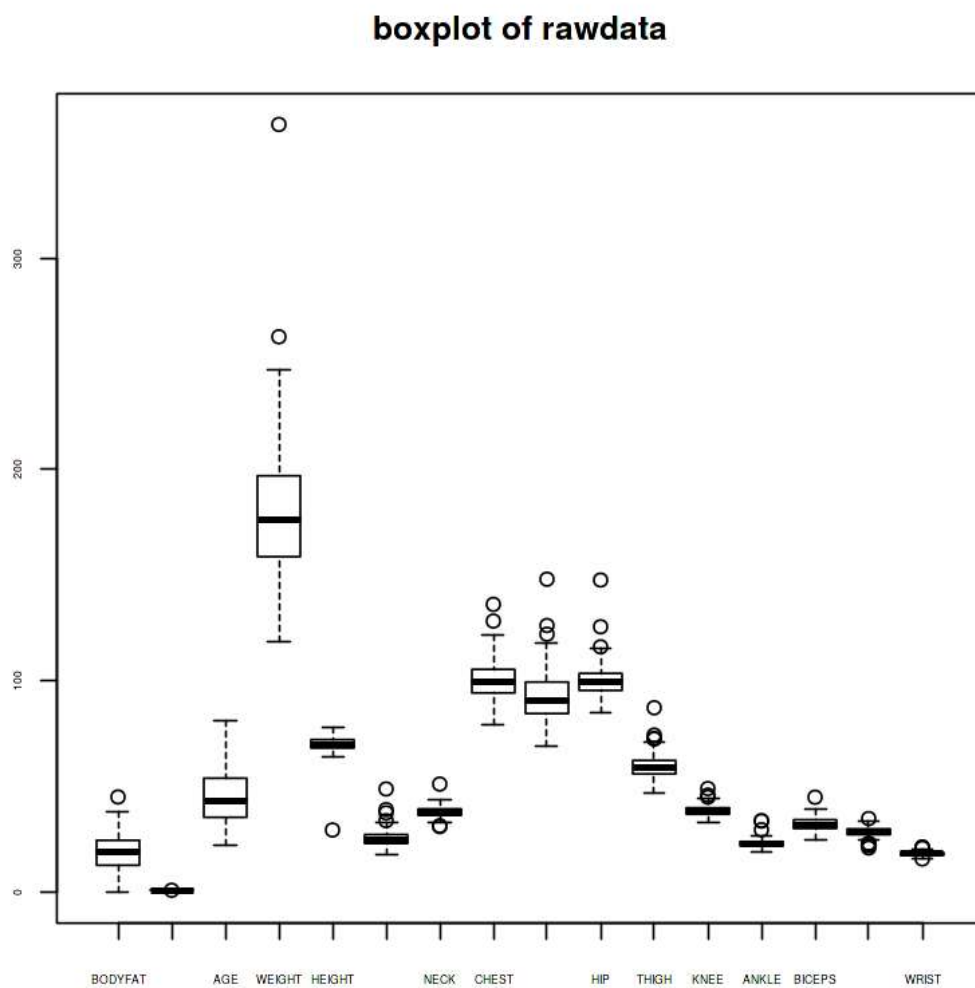
2 Part 1: Dealing with raw data

Read in the raw data and do data cleaning (detect and delete the points which bodyfat does not have a linear relationship of 1 / density)


```
[2]: data = read.csv("BodyFat.csv", row.names = 1)
      boxplot(data,main="boxplot of rawdata",cex.axis=0.4)

      plot(data$BODYFAT,data$DENSITY,xlab="bodyfat",ylab="density")
      text(1:length(data[,1]), data$BODYFAT, 1:length(data[,1]),cex=0.5)
      B = 495*(1/data$DENSITY)-450
      plot(B-data$BODYFAT,xlab = 'Observations', ylab = 'B-Bodyfat(%)', main = "
      ↪"siri's equation")
      text(c(48,76,96,182,216),
      ↪B[c(48,76,96,182,216)]-data$BODYFAT[c(48,76,96,182,216)],
      ↪c(48,76,96,182,216),cex=0.8,pos=4)
```

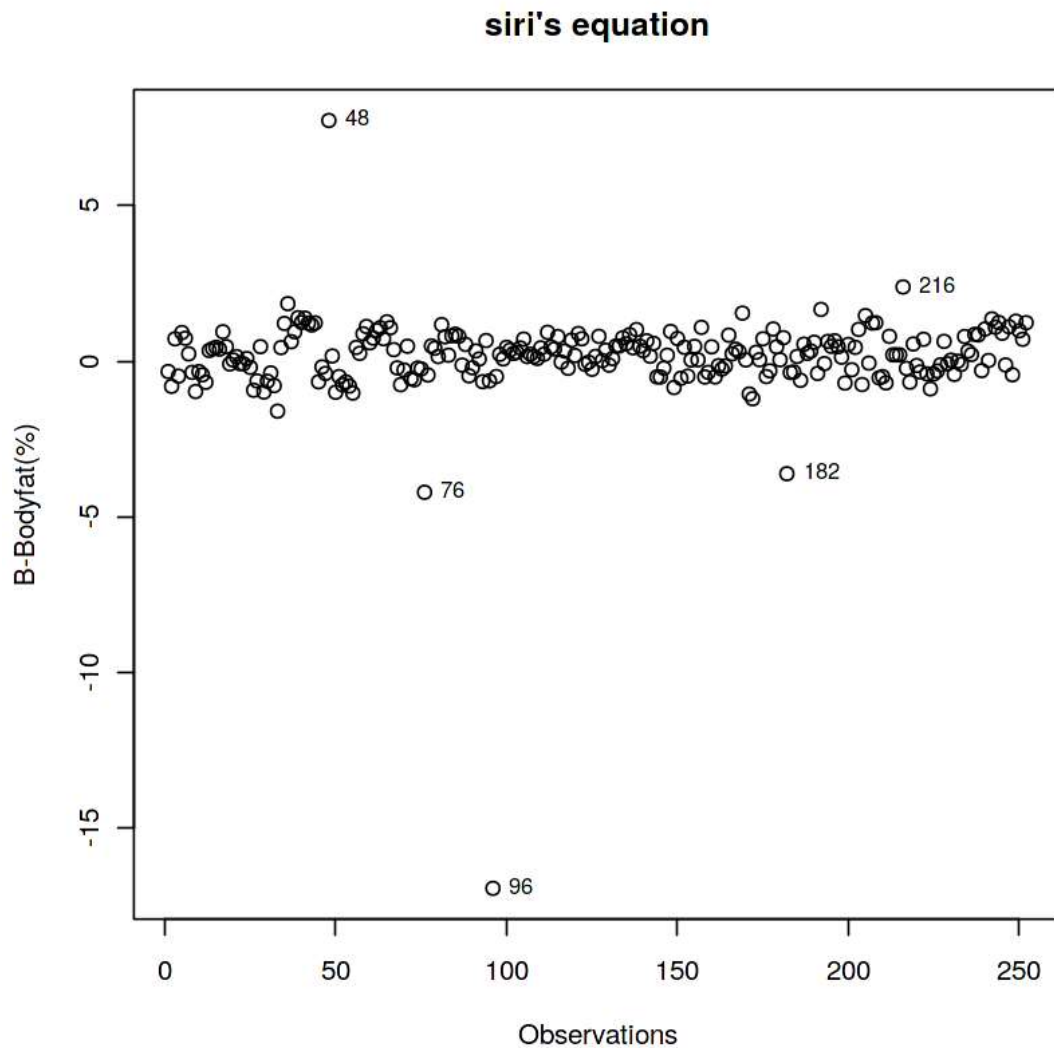
[2]:



[2]:



[2] :




Continue data cleaning (detect and delete the points having wrong relationship of ADIPOSIT, WEIGHT and HEIGHT), then get rid of DENSITY

```
[3]: BMI=703*(data$WEIGHT/(data$HEIGHT)^2)
plot(BMI-data$ADIPOSIT)
text(1:length(data[,1]), BMI-data$ADIPOSIT, 1:length(data[,1]),cex=1)

# get rid of density and get a new dataset
data_new = data[, -2]
```

[3]:



Then delete possible outliers:


```
[4]: #check the cook's distance one by one
model1<- lm(BODYFAT ~ ., data = data_new)
plot(model1, which = 4)
#check the possible outliers
data[c(39, 42, 48, 96,86, 76, 163,182,221),]
```

[4]:

A data.frame: 9 × 16

	BODYFAT <dbl>	DENSITY <dbl>	AGE <int>	WEIGHT <dbl>	HEIGHT <dbl>	ADIPOSITIVITY <dbl>	NECK <dbl>	CHEST <dbl>
39	33.8	1.0202	46	363.15	72.25	48.9	51.2	136.2
42	31.7	1.0250	44	205.00	29.50	29.9	36.6	106.0
48	6.4	1.0665	39	148.50	71.25	20.6	34.6	89.8
96	17.3	1.0991	53	224.50	77.75	26.1	41.1	113.2
86	25.8	1.0386	67	167.00	67.50	26.0	36.5	98.9
76	18.3	1.0666	61	148.25	67.50	22.9	36.0	91.6
163	13.3	1.0690	33	184.25	68.75	24.4	40.7	98.9
182	0.0	1.1089	40	118.50	68.00	18.1	33.8	79.3
221	12.7	1.0706	54	153.25	70.50	24.5	38.5	99.0

[4]:



No.39 has too large weight; No.42 is way too short.

No.48,No.76,No.96 does not match the relationship between bodyfat and density.


No.182 has bodyfat 0, it is a mistake.

No.221,No.163 don't obey the BMI equation (also No.42).

We remove them and scale the data to get a new dataset. Then check the Cook's distance again.

```
[5]: data_clean = data_new[c(-39,-42, -48, -96, -76, -163, -182, -221), ]  
data_clean<-data.frame(scale(data_clean))#scale the data  
write.csv(data_clean,"bodyfat_clean.csv",row.names = F)  
  
#check the cook's distance again  
model = lm(BODYFAT ~ ., data = data_clean)  
plot(model, which = 4)
```

[5]:



We check new data's Cook's distance again and find nothing weird.

3 Part 2: Model selection

The methods we use are: AIC, BIC, Mallow's C_p , Adjusted- R^2 and lasso.

```
[6]: full = lm(BODYFAT ~ ., data = data_clean)
     base = lm(BODYFAT ~ 1, data = data_clean)
```

1. AIC (Backward, forward and both sides)


```
[7]: model_AIC_b = step(full, direction = "backward", trace = 0)
print(model_AIC_b)
summary(model_AIC_b)

model_AIC_f = step(base, direction = "forward", trace = 0, scope = list(lower =
↪base, upper = full))
print(model_AIC_f)

model_AIC_t = step(base, direction = "both", trace = 0, scope = list(lower =
↪base, upper = full))
print(model_AIC_t)
```

Call:

```
lm(formula = BODYFAT ~ AGE + WEIGHT + HEIGHT + ADIPOSITY + NECK +
    ABDOMEN + HIP + THIGH + FOREARM + WRIST, data = data_clean)
```

Coefficients:

(Intercept)	AGE	WEIGHT	HEIGHT	ADIPOSITY	NECK
1.413e-15	9.644e-02	-1.021e+00	3.811e-01	7.602e-01	-9.918e-02
ABDOMEN	HIP	THIGH	FOREARM	WRIST	
1.076e+00	-1.642e-01	1.701e-01	8.393e-02	-1.678e-01	

[7]:

Call:

```
lm(formula = BODYFAT ~ AGE + WEIGHT + HEIGHT + ADIPOSITY + NECK +
    ABDOMEN + HIP + THIGH + FOREARM + WRIST, data = data_clean)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.25248	-0.35972	-0.03786	0.36854	1.20348

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.413e-15	3.276e-02	0.000	1.00000
AGE	9.644e-02	4.809e-02	2.005	0.04608 *
WEIGHT	-1.021e+00	4.902e-01	-2.083	0.03831 *
HEIGHT	3.811e-01	2.435e-01	1.565	0.11896
ADIPOSITY	7.602e-01	4.288e-01	1.773	0.07755 .
NECK	-9.918e-02	6.358e-02	-1.560	0.12017
ABDOMEN	1.076e+00	1.111e-01	9.684	< 2e-16 ***
HIP	-1.642e-01	1.133e-01	-1.449	0.14879
THIGH	1.701e-01	7.905e-02	2.152	0.03241 *
FOREARM	8.393e-02	4.880e-02	1.720	0.08680 .
WRIST	-1.678e-01	5.704e-02	-2.942	0.00359 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5117 on 233 degrees of freedom
Multiple R-squared: 0.7489, Adjusted R-squared: 0.7381
F-statistic: 69.49 on 10 and 233 DF, p-value: < 2.2e-16

Call:

```
lm(formula = BODYFAT ~ ABDOMEN + WEIGHT + WRIST + BICEPS, data = data_clean)
```

Coefficients:

(Intercept)	ABDOMEN	WEIGHT	WRIST	BICEPS
2.369e-16	1.204e+00	-4.139e-01	-1.477e-01	1.077e-01

Call:

```
lm(formula = BODYFAT ~ ABDOMEN + WEIGHT + WRIST + BICEPS, data = data_clean)
```

Coefficients:

(Intercept)	ABDOMEN	WEIGHT	WRIST	BICEPS
2.369e-16	1.204e+00	-4.139e-01	-1.477e-01	1.077e-01

2. BIC (Backward, forward and both sides)

```
[8]: n = dim(data_clean)[1]
model_BIC_b = step(full, direction = "backward", trace = 0, k = log(n))
print(model_BIC_b)
summary(model_BIC_b)

model_BIC_f = step(base, direction = "forward", trace = 0, scope = list(lower = _
→base, upper = full), k = log(n))
print(model_BIC_f )

model_BIC_t = step(full, direction = "both", trace = 0, k = log(n))
print(model_BIC_t)
summary(model_BIC_t)
```

Call:

```
lm(formula = BODYFAT ~ WEIGHT + ABDOMEN + WRIST, data = data_clean)
```

Coefficients:

(Intercept)	WEIGHT	ABDOMEN	WRIST
1.574e-16	-3.238e-01	1.189e+00	-1.376e-01

[8]:

```
Call:
lm(formula = BODYFAT ~ WEIGHT + ABDOMEN + WRIST, data = data_clean)

Residuals:
    Min       1Q   Median       3Q      Max
-1.18875 -0.37533 -0.03202  0.38206  1.20597

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.574e-16  3.329e-02   0.000  1.00000
WEIGHT       -3.238e-01  7.969e-02  -4.063  6.58e-05 ***
ABDOMEN       1.189e+00  6.861e-02  17.329 < 2e-16 ***
WRIST        -1.376e-01  4.878e-02  -2.821  0.00519 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5199 on 240 degrees of freedom
Multiple R-squared:  0.733,    Adjusted R-squared:  0.7297
F-statistic: 219.6 on 3 and 240 DF,  p-value: < 2.2e-16
```

```
Call:
lm(formula = BODYFAT ~ ABDOMEN + WEIGHT + WRIST, data = data_clean)

Coefficients:
(Intercept)      ABDOMEN      WEIGHT      WRIST
  1.574e-16    1.189e+00   -3.238e-01   -1.376e-01
```

```
Call:
lm(formula = BODYFAT ~ WEIGHT + ABDOMEN + WRIST, data = data_clean)

Coefficients:
(Intercept)      WEIGHT      ABDOMEN      WRIST
  1.574e-16   -3.238e-01    1.189e+00   -1.376e-01
```

[8]:

```
Call:
lm(formula = BODYFAT ~ WEIGHT + ABDOMEN + WRIST, data = data_clean)

Residuals:
    Min       1Q   Median       3Q      Max
-1.18875 -0.37533 -0.03202  0.38206  1.20597
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.574e-16	3.329e-02	0.000	1.00000
WEIGHT	-3.238e-01	7.969e-02	-4.063	6.58e-05 ***
ABDOMEN	1.189e+00	6.861e-02	17.329	< 2e-16 ***
WRIST	-1.376e-01	4.878e-02	-2.821	0.00519 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5199 on 240 degrees of freedom
Multiple R-squared: 0.733, Adjusted R-squared: 0.7297
F-statistic: 219.6 on 3 and 240 DF, p-value: < 2.2e-16

3. Mallow's cp

```
[9]: X = data_clean[, -1]
Y = data_clean$BODYFAT
g = leaps(X, Y, nbest = 1)
Cpplot(g)
print(colnames(data_clean)[c(1, 2, 3, 4, 5, 7,8,9,13,14) + 1])
cp_model = BODYFAT ~ AGE +WEIGHT+ HEIGHT +ADIPOSITIY+NECK+ ABDOMEN +HIP+ THIGH_
↪+FOREARM + WRIST
cp_lm = lm(cp_model, data=data_clean)
summary(cp_lm)
```

```
[1] "AGE"      "WEIGHT"   "HEIGHT"   "ADIPOSITIY" "NECK"     "ABDOMEN"
[7] "HIP"      "THIGH"    "FOREARM"  "WRIST"
```

[9]:

Call:

```
lm(formula = cp_model, data = data_clean)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.25248	-0.35972	-0.03786	0.36854	1.20348

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.413e-15	3.276e-02	0.000	1.00000
AGE	9.644e-02	4.809e-02	2.005	0.04608 *
WEIGHT	-1.021e+00	4.902e-01	-2.083	0.03831 *
HEIGHT	3.811e-01	2.435e-01	1.565	0.11896
ADIPOSITIY	7.602e-01	4.288e-01	1.773	0.07755 .
NECK	-9.918e-02	6.358e-02	-1.560	0.12017
ABDOMEN	1.076e+00	1.111e-01	9.684	< 2e-16 ***
HIP	-1.642e-01	1.133e-01	-1.449	0.14879

THIGH	1.701e-01	7.905e-02	2.152	0.03241	*
FOREARM	8.393e-02	4.880e-02	1.720	0.08680	.
WRIST	-1.678e-01	5.704e-02	-2.942	0.00359	**


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5117 on 233 degrees of freedom

Multiple R-squared: 0.7489, Adjusted R-squared: 0.7381

F-statistic: 69.49 on 10 and 233 DF, p-value: < 2.2e-16

[9]:



4. Adjusted R square

```
[10]: g_ad = leaps(X, Y, nbest = 1, method="adjr2")
plot(g_ad$adjr2)
(g_ad$which)[which(g_ad$adjr2==max(g_ad$adjr2)),]
print(colnames(data_clean)[c(1:9, 12:14) + 1])
Ad_r=BODYFAT~AGE +WEIGHT+ HEIGHT +ADIPOSITIVITY+NECK+CHEST+ ABDOMEN +HIP+ THIGH+
↪+BICEPS+FOREARM + WRIST
adr_model<-lm(Ad_r,data=data_clean)
print(adr_model)
```

```
[10]: 1 TRUE 2 TRUE 3 TRUE 4 TRUE 5 TRUE 6 TRUE 7 TRUE 8 TRUE 9 TRUE A FALSE B FALSE
C          TRUE D          TRUE E          TRUE

[1] "AGE"      "WEIGHT"    "HEIGHT"    "ADIPOSITIVITY" "NECK"      "CHEST"
[7] "ABDOMEN"   "HIP"       "THIGH"     "BICEPS"      "FOREARM"   "WRIST"
```


Call:

```
lm(formula = Ad_r, data = data_clean)
```

Coefficients:

(Intercept)	AGE	WEIGHT	HEIGHT	ADIPOSITIVITY	NECK
1.451e-15	9.742e-02	-9.867e-01	3.979e-01	8.468e-01	-1.014e-01
CHEST	ABDOMEN	HIP	THIGH	BICEPS	FOREARM
-1.483e-01	1.123e+00	-1.831e-01	1.208e-01	6.640e-02	6.720e-02
WRIST					
-1.780e-01					

```
[10]:
```



5. Lasso


```
[11]: model_l <- glmnet(as.matrix(data_clean[,2:15]), data_clean$BODYFAT, family = "gaussian",
  ↪ nlambda = 50, alpha = 1, standardize = T)
print(model_l)
plot(model_l, xvar = "lambda", label = T)
cvfit <- cv.glmnet(as.matrix(data_clean[,2:15]), data_clean$BODYFAT, family = "gaussian",
  ↪ type.measure = "mse", nfolds = 10, alpha = 1)
plot(cvfit)
print(coef(cvfit, s = "lambda.1se" ))
model_lasso <- BODYFAT ~ AGE + HEIGHT + ABDOMEN + WRIST
model_l <- lm(BODYFAT ~ AGE + HEIGHT + ABDOMEN + WRIST, data = data_clean)
print(model_l)
```

```
Call: glmnet(x = as.matrix(data_clean[, 2:15]), y = data_clean$BODYFAT,
family = "gaussian", alpha = 1, nlambda = 50, standardize = T)
```

	Df	%Dev	Lambda
[1,]	0	0.0000	0.8218000
[2,]	1	0.2125	0.6809000
[3,]	1	0.3584	0.5643000
[4,]	1	0.4586	0.4676000
[5,]	1	0.5273	0.3875000
[6,]	1	0.5746	0.3211000
[7,]	1	0.6070	0.2660000
[8,]	1	0.6293	0.2205000
[9,]	1	0.6446	0.1827000
[10,]	2	0.6557	0.1514000
[11,]	2	0.6732	0.1254000
[12,]	2	0.6852	0.1039000
[13,]	3	0.6936	0.0861300
[14,]	4	0.7041	0.0713700
[15,]	4	0.7137	0.0591400
[16,]	4	0.7203	0.0490100
[17,]	4	0.7249	0.0406100
[18,]	4	0.7280	0.0336500
[19,]	5	0.7303	0.0278800
[20,]	5	0.7324	0.0231100
[21,]	7	0.7346	0.0191500
[22,]	7	0.7366	0.0158700
[23,]	9	0.7384	0.0131500
[24,]	10	0.7405	0.0108900
[25,]	11	0.7427	0.0090270
[26,]	11	0.7445	0.0074810
[27,]	11	0.7457	0.0061990
[28,]	12	0.7467	0.0051360
[29,]	13	0.7474	0.0042560
[30,]	13	0.7479	0.0035270
[31,]	14	0.7483	0.0029230
[32,]	14	0.7493	0.0024220
[33,]	13	0.7498	0.0020070
[34,]	13	0.7499	0.0016630
[35,]	14	0.7500	0.0013780
[36,]	14	0.7506	0.0011420
[37,]	14	0.7512	0.0009462
[38,]	14	0.7517	0.0007840
[39,]	14	0.7520	0.0006497
[40,]	14	0.7522	0.0005384
[41,]	14	0.7523	0.0004461
[42,]	14	0.7524	0.0003697
[43,]	14	0.7525	0.0003063


```
[44,] 14 0.7525 0.0002538
[45,] 14 0.7526 0.0002103
[46,] 14 0.7526 0.0001743
```

[11]:



15 x 1 sparse Matrix of class "dgCMatrix"

```
      1
(Intercept) -1.056100e-16
AGE          1.174299e-02
WEIGHT       .
HEIGHT      -8.806614e-02
ADIPOSITIVITY .
NECK         .
CHEST        .
ABDOMEN      7.732062e-01
```

```

HIP .
THIGH .
KNEE .
ANKLE .
BICEPS .
FOREARM .
WRIST -1.752753e-02

```


Call:

```
lm(formula = BODYFAT ~ AGE + HEIGHT + ABDOMEN + WRIST, data = data_clean)
```

Coefficients:

(Intercept)	AGE	HEIGHT	ABDOMEN	WRIST
-3.209e-16	7.994e-02	-1.020e-01	9.444e-01	-2.036e-01

[11]:



4 Part3: Other analysis

We compare the models above by their summary table and plots.

```
[12]: summary(model_AIC_b)
summary(model_AIC_f)
summary(model_AIC_t)

summary(model_BIC_b)
summary(model_BIC_f)
summary(model_BIC_t)

summary(cp_lm)
summary(adr_model)
summary(model_l)
```

[12]:

Call:

```
lm(formula = BODYFAT ~ AGE + WEIGHT + HEIGHT + ADIPOSITIVITY + NECK +
    ABDOMEN + HIP + THIGH + FOREARM + WRIST, data = data_clean)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.25248	-0.35972	-0.03786	0.36854	1.20348

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.413e-15	3.276e-02	0.000	1.00000
AGE	9.644e-02	4.809e-02	2.005	0.04608 *
WEIGHT	-1.021e+00	4.902e-01	-2.083	0.03831 *
HEIGHT	3.811e-01	2.435e-01	1.565	0.11896
ADIPOSITIVITY	7.602e-01	4.288e-01	1.773	0.07755 .
NECK	-9.918e-02	6.358e-02	-1.560	0.12017
ABDOMEN	1.076e+00	1.111e-01	9.684	< 2e-16 ***
HIP	-1.642e-01	1.133e-01	-1.449	0.14879
THIGH	1.701e-01	7.905e-02	2.152	0.03241 *
FOREARM	8.393e-02	4.880e-02	1.720	0.08680 .
WRIST	-1.678e-01	5.704e-02	-2.942	0.00359 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5117 on 233 degrees of freedom

Multiple R-squared: 0.7489, Adjusted R-squared: 0.7381

F-statistic: 69.49 on 10 and 233 DF, p-value: < 2.2e-16

[12]:

```
Call:
lm(formula = BODYFAT ~ ABDOMEN + WEIGHT + WRIST + BICEPS, data = data_clean)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.23562	-0.36528	-0.01922	0.39216	1.12071

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.369e-16	3.308e-02	0.000	1.00000
ABDOMEN	1.204e+00	6.860e-02	17.549	< 2e-16 ***
WEIGHT	-4.139e-01	9.123e-02	-4.536	9.07e-06 ***
WRIST	-1.477e-01	4.875e-02	-3.029	0.00272 **
BICEPS	1.077e-01	5.412e-02	1.990	0.04774 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5168 on 239 degrees of freedom

Multiple R-squared: 0.7373, Adjusted R-squared: 0.7329

F-statistic: 167.7 on 4 and 239 DF, p-value: < 2.2e-16

[12]:

```
Call:
lm(formula = BODYFAT ~ ABDOMEN + WEIGHT + WRIST + BICEPS, data = data_clean)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.23562	-0.36528	-0.01922	0.39216	1.12071

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.369e-16	3.308e-02	0.000	1.00000
ABDOMEN	1.204e+00	6.860e-02	17.549	< 2e-16 ***
WEIGHT	-4.139e-01	9.123e-02	-4.536	9.07e-06 ***
WRIST	-1.477e-01	4.875e-02	-3.029	0.00272 **
BICEPS	1.077e-01	5.412e-02	1.990	0.04774 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5168 on 239 degrees of freedom

Multiple R-squared: 0.7373, Adjusted R-squared: 0.7329

F-statistic: 167.7 on 4 and 239 DF, p-value: < 2.2e-16

[12]:

```
Call:
lm(formula = BODYFAT ~ WEIGHT + ABDOMEN + WRIST, data = data_clean)

Residuals:
    Min       1Q   Median       3Q      Max
-1.18875 -0.37533 -0.03202  0.38206  1.20597

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.574e-16  3.329e-02   0.000  1.00000
WEIGHT       -3.238e-01  7.969e-02  -4.063  6.58e-05 ***
ABDOMEN       1.189e+00  6.861e-02  17.329 < 2e-16 ***
WRIST        -1.376e-01  4.878e-02  -2.821  0.00519 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5199 on 240 degrees of freedom
Multiple R-squared:  0.733,    Adjusted R-squared:  0.7297
F-statistic: 219.6 on 3 and 240 DF,  p-value: < 2.2e-16
```

[12]:

```
Call:
lm(formula = BODYFAT ~ ABDOMEN + WEIGHT + WRIST, data = data_clean)

Residuals:
    Min       1Q   Median       3Q      Max
-1.18875 -0.37533 -0.03202  0.38206  1.20597

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.574e-16  3.329e-02   0.000  1.00000
ABDOMEN       1.189e+00  6.861e-02  17.329 < 2e-16 ***
WEIGHT       -3.238e-01  7.969e-02  -4.063  6.58e-05 ***
WRIST        -1.376e-01  4.878e-02  -2.821  0.00519 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5199 on 240 degrees of freedom
Multiple R-squared:  0.733,    Adjusted R-squared:  0.7297
F-statistic: 219.6 on 3 and 240 DF,  p-value: < 2.2e-16
```

[12]:

```
Call:
lm(formula = BODYFAT ~ WEIGHT + ABDOMEN + WRIST, data = data_clean)

Residuals:
    Min       1Q   Median       3Q      Max
-1.18875 -0.37533 -0.03202  0.38206  1.20597

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.574e-16  3.329e-02   0.000  1.00000
WEIGHT       -3.238e-01  7.969e-02  -4.063  6.58e-05 ***
ABDOMEN       1.189e+00  6.861e-02  17.329 < 2e-16 ***
WRIST        -1.376e-01  4.878e-02  -2.821  0.00519 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5199 on 240 degrees of freedom
Multiple R-squared:  0.733,    Adjusted R-squared:  0.7297
F-statistic: 219.6 on 3 and 240 DF,  p-value: < 2.2e-16
```

[12]:

```
Call:
lm(formula = cp_model, data = data_clean)

Residuals:
    Min       1Q   Median       3Q      Max
-1.25248 -0.35972 -0.03786  0.36854  1.20348

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  1.413e-15  3.276e-02   0.000  1.00000
AGE          9.644e-02  4.809e-02   2.005  0.04608 *
WEIGHT      -1.021e+00  4.902e-01  -2.083  0.03831 *
HEIGHT       3.811e-01  2.435e-01   1.565  0.11896
ADIPOSIITY   7.602e-01  4.288e-01   1.773  0.07755 .
NECK        -9.918e-02  6.358e-02  -1.560  0.12017
ABDOMEN      1.076e+00  1.111e-01   9.684 < 2e-16 ***
HIP         -1.642e-01  1.133e-01  -1.449  0.14879
THIGH        1.701e-01  7.905e-02   2.152  0.03241 *
FOREARM      8.393e-02  4.880e-02   1.720  0.08680 .
WRIST       -1.678e-01  5.704e-02  -2.942  0.00359 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5117 on 233 degrees of freedom
```

Multiple R-squared: 0.7489, Adjusted R-squared: 0.7381
F-statistic: 69.49 on 10 and 233 DF, p-value: < 2.2e-16

[12]:

Call:

```
lm(formula = Ad_r, data = data_clean)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.24761	-0.34427	-0.02745	0.37215	1.16835

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.451e-15	3.270e-02	0.000	1.00000
AGE	9.742e-02	4.824e-02	2.019	0.04460 *
WEIGHT	-9.867e-01	4.954e-01	-1.992	0.04757 *
HEIGHT	3.979e-01	2.433e-01	1.635	0.10338
ADIPOSIT	8.468e-01	4.323e-01	1.959	0.05134 .
NECK	-1.014e-01	6.367e-02	-1.592	0.11273
CHEST	-1.483e-01	1.086e-01	-1.365	0.17357
ABDOMEN	1.123e+00	1.143e-01	9.825	< 2e-16 ***
HIP	-1.831e-01	1.155e-01	-1.585	0.11440
THIGH	1.208e-01	8.401e-02	1.438	0.15169
BICEPS	6.640e-02	6.020e-02	1.103	0.27115
FOREARM	6.720e-02	5.070e-02	1.326	0.18631
WRIST	-1.780e-01	5.731e-02	-3.107	0.00213 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5107 on 231 degrees of freedom

Multiple R-squared: 0.752, Adjusted R-squared: 0.7391

F-statistic: 58.38 on 12 and 231 DF, p-value: < 2.2e-16

[12]:

Call:

```
lm(formula = BODYFAT ~ AGE + HEIGHT + ABDOMEN + WRIST, data = data_clean)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.25508	-0.37499	-0.03633	0.39655	1.14166

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-3.209e-16	3.324e-02	0.000	1.00000
AGE	7.994e-02	3.722e-02	2.148	0.03272 *

HEIGHT	-1.020e-01	3.905e-02	-2.612	0.00958	**
ABDOMEN	9.444e-01	4.223e-02	22.361	< 2e-16	***
WRIST	-2.036e-01	4.591e-02	-4.434	1.41e-05	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5193 on 239 degrees of freedom



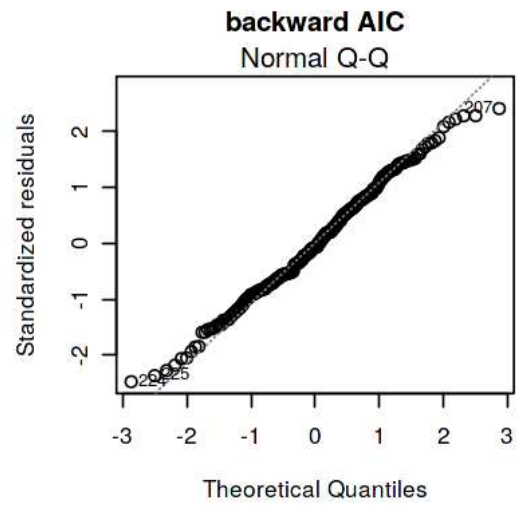

Multiple R-squared: 0.7348, Adjusted R-squared: 0.7304

F-statistic: 165.5 on 4 and 239 DF, p-value: < 2.2e-16




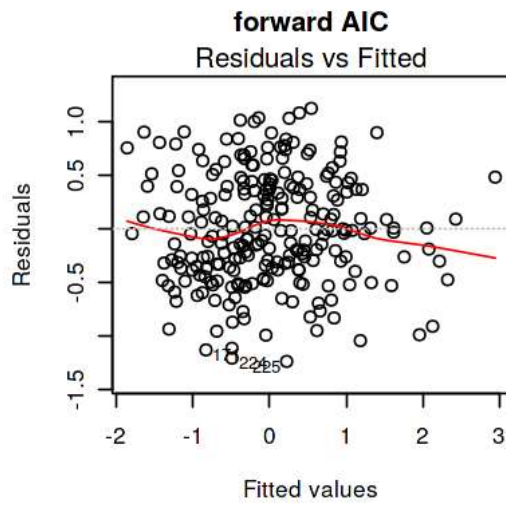
4.1 Diagnostic plot

```
[13]: layout(matrix(1:4, byrow = TRUE, nrow = 2))
      plot(model_AIC_b,main="backward AIC")
      plot(model_AIC_f,main="forward AIC")
      plot(model_AIC_t,main="both sided AIC")
      plot(model_BIC_b,main="backward BIC")
      plot(model_BIC_f,main="forward BIC")
      plot(model_BIC_t,main="both sided BIC")
      plot(model_l,main="lasso model")
      plot(cp_lm,main="Mallow's cp model")
      plot(adr_model,main="adjusted Rsq model")
```

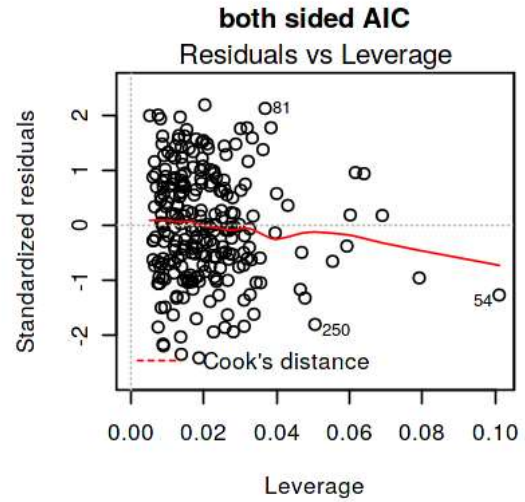

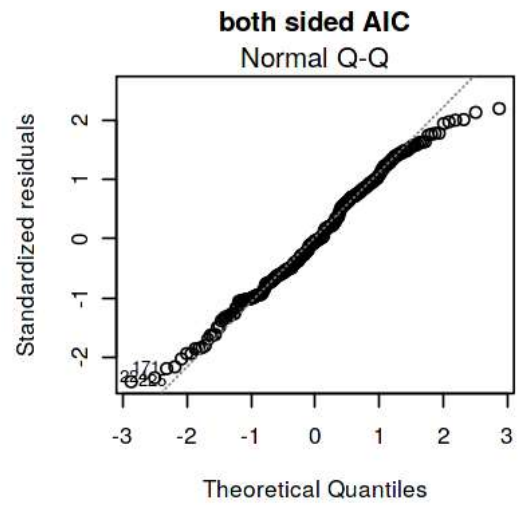

[13]:





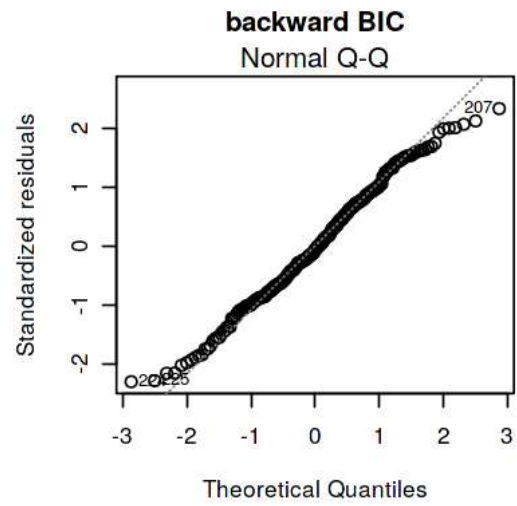
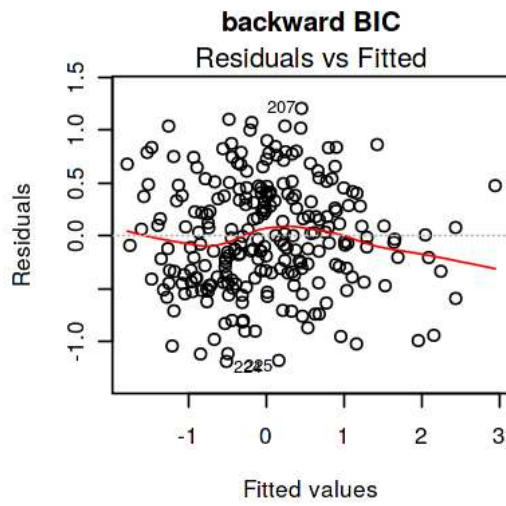
[13]:



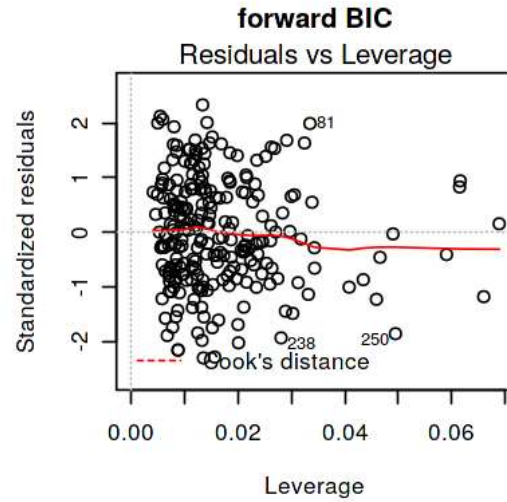



[13] :




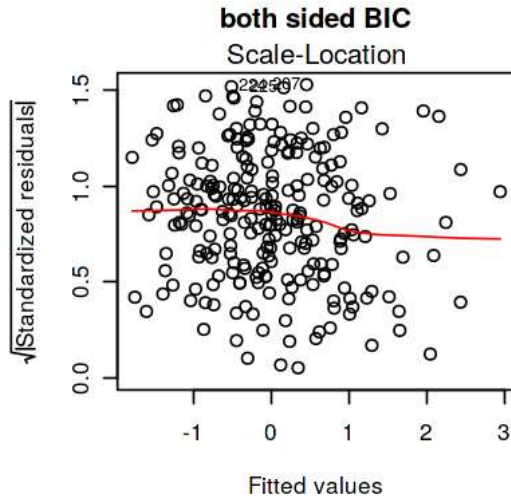


[13] :







[13]:







[13] :



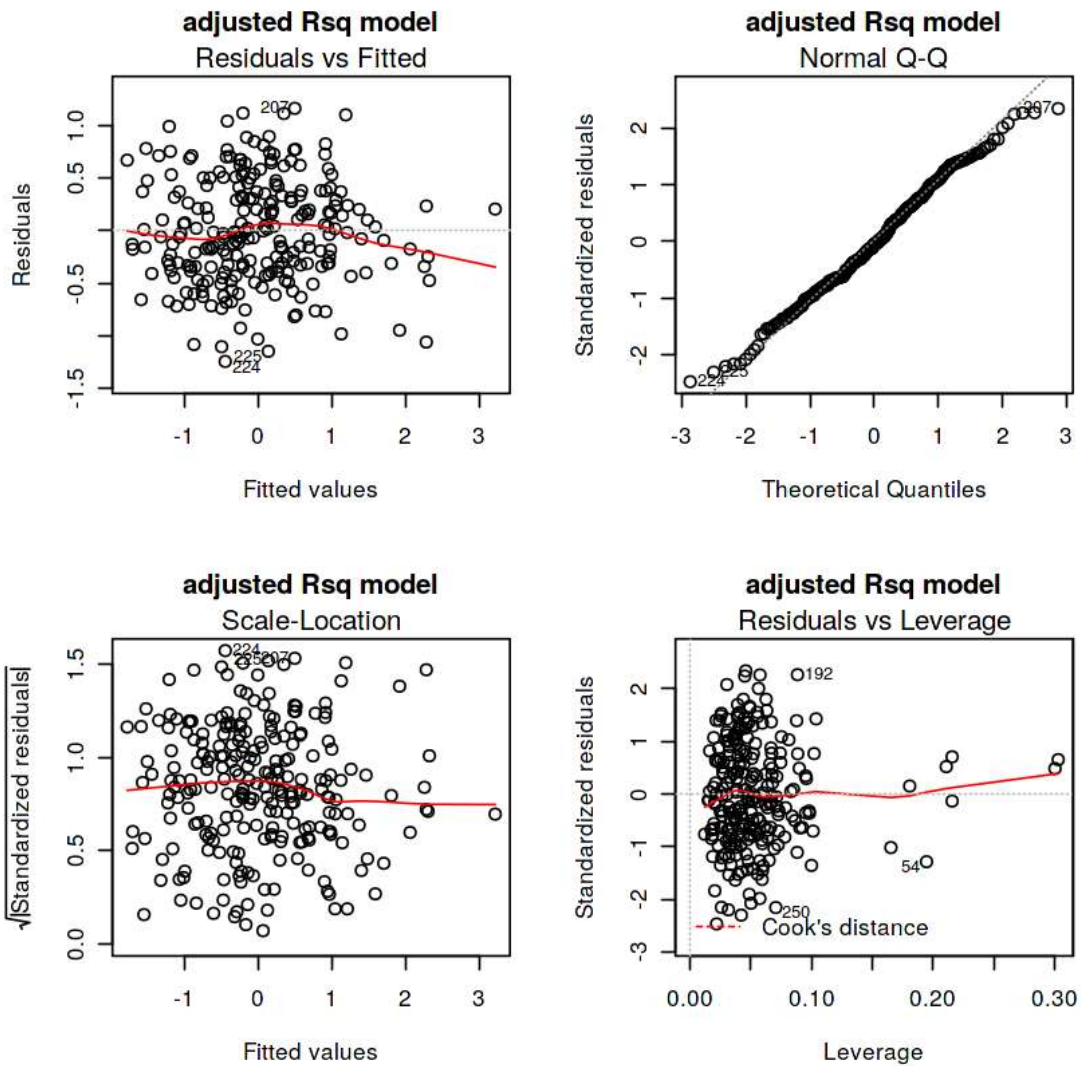
[13] :



[13] :



[13]:




Also do ANOVA table among them, model BIC forward is the best considering the F statistics and adjusted R square.

[0]:

```
[14]: plot(model_BIC_f$residuals,main="Residual Plot")
      abline(h=0,col="red")
```

[14]:



4.1.1 According to the model selection steps above, it seems that *abdomen*, *weight* and *wrist* are three most important predictors. By “rule of thumb” we can choose our predictors among them. Now check the possible multicollinearity between abdomen, wrist and weight. We create cleaned data that are not scaled again

```
[15]: new_data_clean = data_new[c(-39, -42, -48, -96, -76, -163, -182, -221), ]
```

```
[16]: cor(new_data_clean$WRIST,new_data_clean$ABDOMEN)
      cor(new_data_clean$WEIGHT,new_data_clean$ABDOMEN)
      cor(new_data_clean$WRIST,new_data_clean$WEIGHT)
```

```
vif(lm(BODYFAT ~ ABDOMEN + WEIGHT + WRIST, data = new_data_clean))
vif(lm(BODYFAT ~ ABDOMEN + WEIGHT, data = new_data_clean))
vif(lm(BODYFAT ~ ABDOMEN + WRIST, data = new_data_clean))
```

```
[16]: 0.603049909542607
```

```
[16]: 0.872690169082953
```

```
[16]: 0.726881563271104
```

```
[16]: ABDOMEN          4.23126966566092 WEIGHT          5.70873762070487 WRIST
      2.13887303654502
```

```
[16]: ABDOMEN          4.19442205244338 WEIGHT          4.19442205244298
```

```
[16]: ABDOMEN          1.57150964502599 WRIST          1.57150964502593
```

It seems that these predictors are correlated but can not be considered as multicollinearity. We consider using only two variables and see the model's performance.

We consider using only two variables or even only one to see the model's performance-try to use only two x and check the model results

```
[17]: summary(lm(BODYFAT ~ ABDOMEN + WRIST, data = new_data_clean))
summary(lm(BODYFAT ~ ABDOMEN + WEIGHT, data = new_data_clean))
```

```
[17]:
```

Call:

```
lm(formula = BODYFAT ~ ABDOMEN + WRIST, data = new_data_clean)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-9.3973	-2.9225	-0.2025	2.9356	10.3555

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-11.19959	5.39785	-2.075	0.0391 *
ABDOMEN	0.72574	0.03234	22.440	< 2e-16 ***
WRIST	-2.02533	0.36453	-5.556	7.29e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.087 on 241 degrees of freedom

Multiple R-squared: 0.7146, Adjusted R-squared: 0.7123

F-statistic: 301.8 on 2 and 241 DF, p-value: < 2.2e-16

```
[17]:
```

Call:

```
lm(formula = BODYFAT ~ ABDOMEN + WEIGHT, data = new_data_clean)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-10.129	-2.897	-0.116	2.880	9.816

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-42.30901	2.45220	-17.253	< 2e-16 ***
ABDOMEN	0.90497	0.05195	17.420	< 2e-16 ***
WEIGHT	-0.12520	0.01974	-6.344	1.1e-09 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.018 on 241 degrees of freedom

Multiple R-squared: 0.7241, Adjusted R-squared: 0.7218


F-statistic: 316.3 on 2 and 241 DF, p-value: < 2.2e-16

By summary table, we can see the R-square value of “ABDOMEN and WRIST” model and ABDOMEN and WEIGHT are almost the same. Also, by VIF analysis we find the VIF of “ABDOMEN and WEIGHT” model is larger than “ABDOMEN and WRIST” model.


So our final model should be $\text{BODYFAT} \sim \text{ABDOMEN} + \text{WRIST}$ because it is they are two most important predictors and it can be easily measured. Then we do model diagnostics.

```
[18]: final_model<-lm(BODYFAT ~ ABDOMEN + WRIST, data = new_data_clean)
plot(final_model,which=1)
plot(final_model,which=2)
plot(final_model,which=4)
```


[18]:



[18]:



[18] :



Advantages of Our Model

1. Our final model is simple and straightforward. The variables included for prediction is easy to get.

2. The model is robust.

3.The model is accurate as the adjusted R-squared is greater than 0.6.